

Multiple file linking in C: A short detour

Ack: Example from
slides on [Linking](#)

15-213: Introduction to Computer Systems
11th Lecture, Sept. 30, 2010

Instructors:
Randy Bryant and Dave O'Hallaron

1

Multiple file linking in C: A short detour

```
/* main.c */
#include <stdio.h>
void swap();
int buf[2] = {0x137, 0x291};
int main()
{
    printf("%d, %d\n", buf[0], buf[1]);
    swap();
    printf("%d, %d\n", buf[0], buf[1]);
    return 0;
}
```

Ack: Example from
slides on [Linking](#)

15-213: Introduction to Computer Systems
11th Lecture, Sept. 30, 2010

Instructors:
Randy Bryant and Dave O'Hallaron

1

Multiple file linking in C: A short detour

```
/* main.c */
#include <stdio.h>
void swap();
int buf[2] = {0x137, 0x291};
int main()
{
    printf("%d, %d\n", buf[0], buf[1]);
    swap();
    printf("%d, %d\n", buf[0], buf[1]);
    return 0;
}
```

Ack: Example from
slides on [Linking](#)

15-213: Introduction to Computer Systems
11th Lecture, Sept. 30, 2010

Instructors:
Randy Bryant and Dave O'Hallaron

1

```
/* swap.c */
extern int buf[];
int *bufp0 = &buf[0];
int *bufp1;
#define BADVALUE 0x999
void swap()
{
    int temp = BADVALUE;
    bufp1 = &buf[1];
    temp = *bufp0;
    *bufp0 = *bufp1;
    *bufp1 = temp;
}
```

How to get an executable program from multiple C files?

4

gcc options

5

gcc options

- Preprocessing only

```
gcc -E main.c
```

-- Output on stdout

5

gcc options

- Preprocessing only

```
gcc -E main.c
```

-- Output on stdout

- Object code generation

```
gcc -c main.c
```

-- Generates main.o

5

gcc options

- Preprocessing only

```
gcc -E main.c
```

-- Output on stdout

- Object code generation

```
gcc -c main.c
```

-- Generates main.o

- Assembly code generation

```
gcc -S main.c
```

-- Generates main.s

5

gcc options

- Preprocessing only
`gcc -E main.c` -- Output on stdout
- Object code generation
`gcc -c main.c` -- Generates main.o
- Assembly code generation
`gcc -S main.c`
- Full compilation only
`gcc main.c swap.c` -- Generates main.s
 -- Generates a.out

5

gcc options

- Use `-g` option to enable debugging

6

objdump

- Usage: `objdump <option(s)> <file(s)>`
- Display information from object <file(s)>

7

objdump

- Usage: `objdump <option(s)> <file(s)>`
- Display information from object <file(s)>
`objdump -d a.out` -- dump only .text section

7

objdump

- Usage: `objdump <option(s)> <file(s)>`
- Display information from object <file(s)>
 - `objdump -d a.out` *-- dump only .text section*
 - `objdump -D a.out` *-- dump all sections*

7

objdump

- Usage: `objdump <option(s)> <file(s)>`
- Display information from object <file(s)>
 - `objdump -d a.out` *-- dump only .text section*
 - `objdump -D a.out` *-- dump all sections*
 - `objdump -S swap.o` *-- If .o is created with -g display source statements*

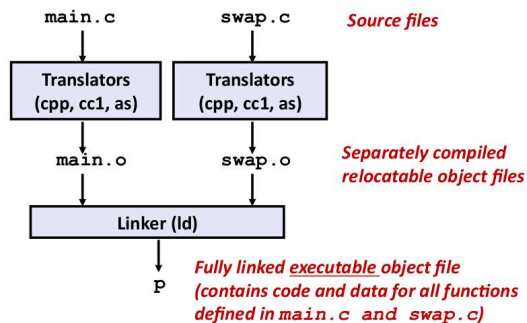
7

Static Linking

- Programs are translated and linked using a *compiler driver*:

```

■ unix> gcc -O2 -g -o p main.c swap.c
■ unix> ./p
  
```



Ack: Bryant & O'Hallaron

Multiple File Linking : WHY?

- Modularity
 - How?
- Efficiency
 - How?

7

What Do Linkers Do?

■ Step 1. Symbol resolution

- Programs define and reference *symbols* (variables and functions):
 - `void swap() {...} /* define symbol swap */`
 - `swap(); /* reference symbol a */`
 - `int *xp = &x; /* define xp, reference x */`
- Symbol definitions are stored (by compiler) in *symbol table*.
 - Symbol table is an array of structs
 - Each entry includes name, size, and location of symbol.
- Linker associates each symbol reference with exactly one symbol definition.

What Do Linkers Do? (cont)

■ Step 2. Relocation

- Merges separate code and data sections into single sections
- Relocates symbols from their relative locations in the `.o` files to their final absolute memory locations in the executable.
- Updates all references to these symbols to reflect their new positions.

Three Kinds of Object Files (Modules)

- **Relocatable object file (.o file)**
 - Contains code and data in a form that can be combined with other relocatable object files to form executable object file.
 - Each `.o` file is produced from exactly one `.c` source
- **Executable object file (a.out file)**
 - Contains code and data in a form that can be copied directly into memory and then executed.
- **Shared object file (.so file)**
 - Special type of relocatable object file that can be loaded into memory and linked dynamically, at either load time or run-time.
 - Called *Dynamic Link Libraries* (DLLs) by Windows

Executable and Linkable Format (ELF)

- **Standard binary format for object files**
- **Originally proposed by AT&T System V Unix**
 - Later adopted by BSD Unix variants and Linux
- **One unified format for**
 - Relocatable object files (`.o`),
 - Executable object files (`a.out`)
 - Shared object files (`.so`)
- **Generic name: ELF binaries**

ELF Object File Format

■ Elf header

- Word size, byte ordering, file type (.o, exec, .so), machine type, etc.

■ Segment header table

- Page size, virtual addresses memory segments (sections), segment sizes.

■ .text section

- Code

■ .rodata section

- Read only data: jump tables, ...

■ .data section

- Initialized global variables

■ .bss section

- Uninitialized global variables
- "Block Started by Symbol"
- "Better Save Space"
- Has section header but occupies no space

ELF header
Segment header table (required for executables)
.text section
.rodata section
.data section
.bss section
.symtab section
.rel.txt section
.rel.data section
.debug section
Section header table

ELF Object File Format (cont.)

■ .symtab section

- Symbol table
- Procedure and static variable names
- Section names and locations

■ .rel.text section

- Relocation info for .text section
- Addresses of instructions that will need to be modified in the executable
- Instructions for modifying.

■ .rel.data section

- Relocation info for .data section
- Addresses of pointer data that will need to be modified in the merged executable

■ .debug section

- Info for symbolic debugging (gcc -g)

■ Section header table

- Offsets and sizes of each section

ELF header
Segment header table (required for executables)
.text section
.rodata section
.data section
.bss section
.symtab section
.rel.txt section
.rel.data section
.debug section
Section header table

Linker Symbols

■ Global symbols

- Symbols defined by module *m* that can be referenced by other modules.
- E.g.: non-**static** C functions and non-**static** global variables.

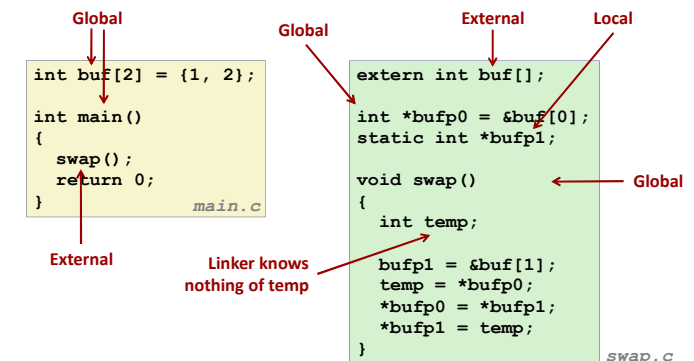
■ External symbols

- Global symbols that are referenced by module *m* but defined by some other module.

■ Local symbols




- Symbols that are defined and referenced exclusively by module *m*.
- E.g.: C functions and variables defined with the **static** attribute.
- Local linker symbols are not local program variables**

Resolving Symbols

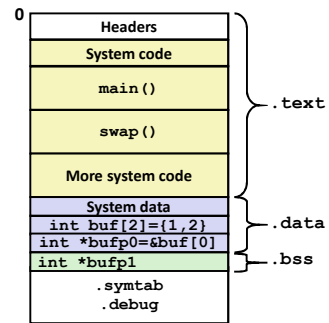


Relocating Code and Data

Relocatable Object Files



Executable Object File



Even though private to swap, requires allocation in .bss

Relocation Info (main)

```

main.c
int buf[2] = {1,2};

int main()
{
    swap();
    return 0;
}

main.o
00000000 <main>:
0: 8d 4c 24 04    lea    0x4(%esp),%ecx
4: 83 e4 f0      and    $0xffffffff0,%esp
7: ff 71 fc      pushl  0xffffffffc(%ecx)
a: 55           push   %ebp
b: 89 e5        mov     %esp,%ebp
d: 51           push   %ecx
e: 83 ec 04     sub     $0x4,%esp
11: e8 fc ff ff   call   12 <main+0x12>
12: R_386_PC32 swap
16: 83 c4 04     add     $0x4,%esp
19: 31 c0        xor     %eax,%eax
1b: 59           pop     %ecx
1c: 5d           pop     %ebp
1d: 8d 61 fc     lea     0xffffffffc(%ecx),%esp
20: c3           ret

```

Disassembly of section .data:

Source: objdump -r -d

```

00000000 <buf>:
0: 01 00 00 00 02 00 00 00

```

Relocation Info (swap, .text)

```

swap.c
extern int buf[];

int
*bufp0 = &buf[0];

static int *bufp1;

void swap()
{
    int temp;

    bufp1 = &buf[1];
    temp = *bufp0;
    *bufp0 = *bufp1;
    *bufp1 = temp;
}

swap.o
Disassembly of section .text:
00000000 <swap>:
0: 8b 15 00 00 00 00    mov     0x0,%edx
2: R_386_32 buf
6: a1 04 00 00 00 00    mov     0x4,%eax
7: R_386_32 buf
b: 55                 push    %ebp
c: 89 e5             mov     %esp,%ebp
e: c7 05 00 00 00 04    movl    $0x4,0x0
15: 00 00 00
10: R_386_32 .bss
14: R_386_32 buf
18: 8b 08             mov     (%eax),%ecx
1a: 89 10             mov     %edx,(%eax)
1c: 5d                 pop     %ebp
1d: 89 0d 04 00 00 00    mov     %ecx,0x4
1f: R_386_32 buf
23: c3                 ret

```

Relocation Info (swap, .data)

```

swap.c
extern int buf[];

int *bufp0 = &buf[0];
static int *bufp1;

void swap()
{
    int temp;

    bufp1 = &buf[1];
    temp = *bufp0;
    *bufp0 = *bufp1;
    *bufp1 = temp;
}

Disassembly of section .data:
00000000 <bufp0>:
0: 00 00 00 00
0: R_386_32 buf

```

Executable Before/After Relocation (.text)

```

00000000 <main>:
    .: 83 ec 04      sub    $0x4,%esp
11: e8 fc ff ff     call   12 <main+0x12>
                                12: R_386_PC32 swap
16: 83 c4 04      add    $0x4,%esp
    .:

08048380 <main>:
08048380: 8d 4c 24 04      lea    0x4(%esp),%ecx
08048384: 83 e4 f0      and    $0xffffffff0,%esp
08048387: ff 71 fc      pushl  0xffffffffc(%ecx)
0804838a: 55          push   %ebp
0804838b: 89 e5      mov    %esp,%ebp
0804838d: 51          push   %ecx
0804838e: 83 ec 04      sub    $0x4,%esp
08048391: e8 1a 00 00 00 call   80483b0 <swap>
08048396: 83 c4 04      add    $0x4,%esp
08048399: 31 c0      xor    %eax,%eax
0804839b: 59          pop     %ecx
0804839c: 5d          pop     %ebp
0804839d: 8d 61 fc      lea    0xffffffffc(%ecx),%esp
080483a0: c3          ret

```

0x8048396 + 0x1a
= 0x80483b0

```

0: 8b 15 00 00 00      mov    0x0,%edx
                                2: R_386_32 buf
6: a1 04 00 00 00      mov    0x4,%eax
                                7: R_386_32 buf
...
e: c7 05 00 00 00 04 movl    $0x4,0x0
15: 00 00 00
                                10: R_386_32 .bss
                                14: R_386_32 buf
...
1d: 89 0d 04 00 00 00      mov    %ecx,0x4
                                1f: R_386_32 buf
23: c3          ret

```

```

080483b0 <swap>:
080483b0: 8b 15 20 96 04 08      mov    0x8049620,%edx
080483b6: a1 24 96 04 08      mov    0x8049624,%eax
080483bb: 55          push   %ebp
080483bc: 89 e5      mov    %esp,%ebp
080483be: c7 05 30 96 04 08 24 movl    $0x8049624,0x8049630
080483c5: 96 04 08
080483c8: 8b 08      mov    (%eax),%ecx
080483ca: 89 10      mov    %edx,(%eax)
080483cc: 5d          pop     %ebp
080483cd: 89 0d 24 96 04 08      mov    %ecx,0x8049624
080483d3: c3          ret

```

Executable After Relocation (.data)

Disassembly of section .data:

```

08049620 <buf>:
8049620: 01 00 00 00 02 00 00 00

08049628 <bufp0>:
8049628: 20 96 04 08

```

Strong and Weak Symbols

- Program symbols are either strong or weak
 - Strong**: procedures and initialized globals
 - Weak**: uninitialized globals

	p1.c	p2.c	
strong	int foo=5;	int foo;	weak
strong	p1() { }	p2() { }	strong

Linker's Symbol Rules

- **Rule 1: Multiple strong symbols are not allowed**
 - Each item can be defined only once
 - Otherwise: Linker error
- **Rule 2: Given a strong symbol and multiple weak symbol, choose the strong symbol**
 - References to the weak symbol resolve to the strong symbol
- **Rule 3: If there are multiple weak symbols, pick an arbitrary one**
 - Can override this with `gcc -fno-common`

Linker Puzzles

<code>int x; p1() {}</code>	<code>p1() {}</code>	Link time error: two strong symbols (p1)
<code>int x; p1() {}</code>	<code>int x; p2() {}</code>	References to <code>x</code> will refer to the same uninitialized int. Is this what you really want?
<code>int x; int y; p1() {}</code>	<code>double x; p2() {}</code>	Writes to <code>x</code> in <code>p2</code> might overwrite <code>y</code> ! Evil!
<code>int x=7; int y=5; p1() {}</code>	<code>double x; p2() {}</code>	Writes to <code>x</code> in <code>p2</code> will overwrite <code>y</code> ! Nasty!
<code>int x=7; p1() {}</code>	<code>int x; p2() {}</code>	References to <code>x</code> will refer to the same initialized variable.

Nightmare scenario: two identical weak structs, compiled by different compilers with different alignment rules.

Role of .h Files

c1.c

```
#include "global.h"

int f() {
    return g+1;
}
```

global.h

```
#ifndef INITIALIZE
int g = 23;
static int init = 1;
#else
int g;
static int init = 0;
#endif
```

c2.c

```
#include <stdio.h>
#include "global.h"

int main() {
    if (!init)
        g = 37;
    int t = f();
    printf("Calling f yields %d\n", t);
    return 0;
}
```

Running Preprocessor

c1.c

```
#include "global.h"

int f() {
    return g+1;
}
```

global.h

```
#ifndef INITIALIZE
int g = 23;
static int init = 1;
#else
int g;
static int init = 0;
#endif
```

-DINITIALIZE

```
int g = 23;
static int init = 1;
int f() {
    return g+1;
}
```

no initialization

```
int g;
static int init = 0;
int f() {
    return g+1;
}
```

#include causes C preprocessor to insert file verbatim

Role of .h Files (cont.)

c1.c

```
#include "global.h"

int f() {
    return g+1;
}
```

global.h

```
#ifndef INITIALIZE
int g = 23;
static int init = 1;
#else
int g;
static int init = 0;
#endif
```

c2.c

```
#include <stdio.h>
#include "global.h"

int main() {
    if (!init)
        g = 37;
    int t = f();
    printf("Calling f yields %d\n", t);
    return 0;
}
```

What happens:

```
gcc -o p c1.c c2.c
??
gcc -o p c1.c c2.c \
-DINITIALIZE
??
```

Global Variables

- **Avoid if you can**
- **Otherwise**
 - Use **static** if you can
 - Initialize if you define a global variable
 - Use **extern** if you use external global variable

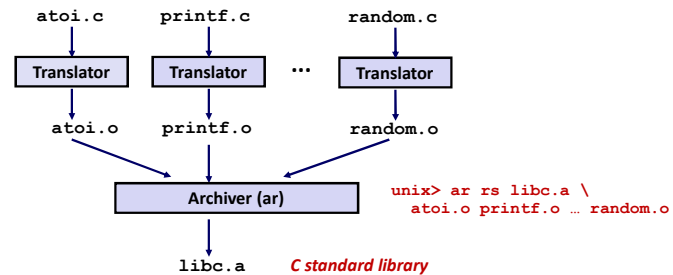
Packaging Commonly Used Functions

- **How to package functions commonly used by programmers?**
 - Math, I/O, memory management, string manipulation, etc.
- **Awkward, given the linker framework so far:**
 - **Option 1:** Put all functions into a single source file
 - Programmers link big object file into their programs
 - Space and time inefficient
 - **Option 2:** Put each function in a separate source file
 - Programmers explicitly link appropriate binaries into their programs
 - More efficient, but burdensome on the programmer

Solution: Static Libraries

- **Static libraries (.a archive files)**
 - Concatenate related relocatable object files into a single file with an index (called an *archive*).
 - Enhance linker so that it tries to resolve unresolved external references by looking for the symbols in one or more archives.
 - If an archive member file resolves reference, link it into the executable.

Creating Static Libraries



- Archiver allows incremental updates
- Recompiles function that changes and replace .o file in archive.

Commonly Used Libraries

libc.a (the C standard library)

- 8 MB archive of 1392 object files.
- I/O, memory allocation, signal handling, string handling, data and time, random numbers, integer math

libm.a (the C math library)

- 1 MB archive of 401 object files.
- floating point math (sin, cos, tan, log, exp, sqrt, ...)

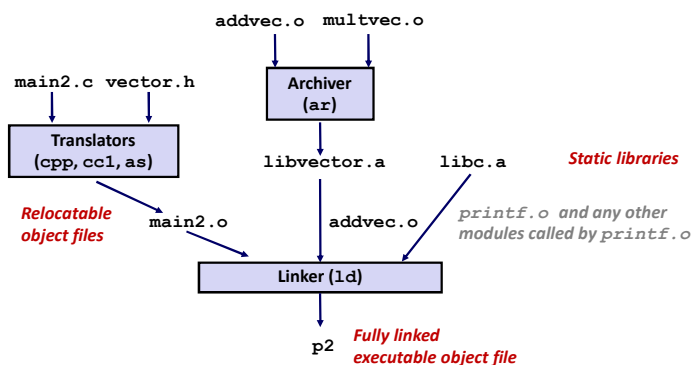
```
% ar -t /usr/lib/libc.a | sort
```

```
...
fork.o
...
fprintf.o
fpu_control.o
fputc.o
freopen.o
fscanf.o
fseek.o
fstab.o
...
```

```
% ar -t /usr/lib/libm.a | sort
```

```
...
_e_acos.o
_e_acosf.o
_e_acosh.o
_e_acoshf.o
_e_acoshl.o
_e_acosl.o
_e_asin.o
_e_asinf.o
_e_asinl.o
...
```

Linking with Static Libraries



Using Static Libraries

■ Linker's algorithm for resolving external references:

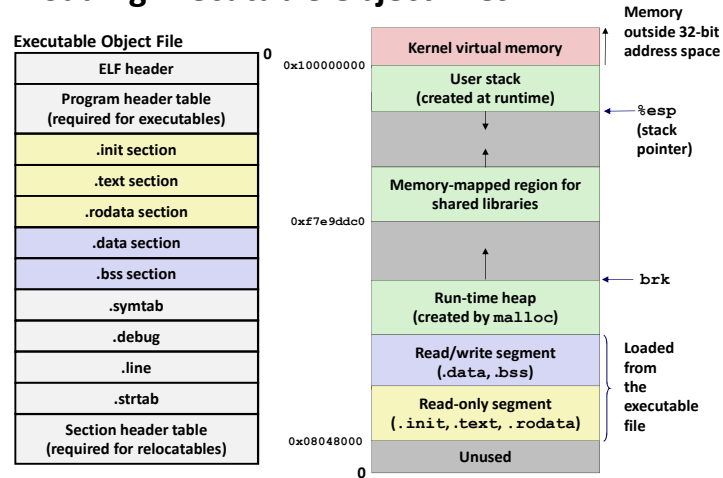
- Scan .o files and .a files in the command line order.
- During the scan, keep a list of the current unresolved references.
- As each new .o or .a file, *obj*, is encountered, try to resolve each unresolved reference in the list against the symbols defined in *obj*.
- If any entries in the unresolved list at end of scan, then error.

■ Problem:

- Command line order matters!
- Moral: put libraries at the end of the command line.

```
unix> gcc -L. libtest.o -lmine
unix> gcc -L. -lmine libtest.o
libtest.o: In function `main':
libtest.o(.text+0x4): undefined reference to `libfun'
```

Loading Executable Object Files



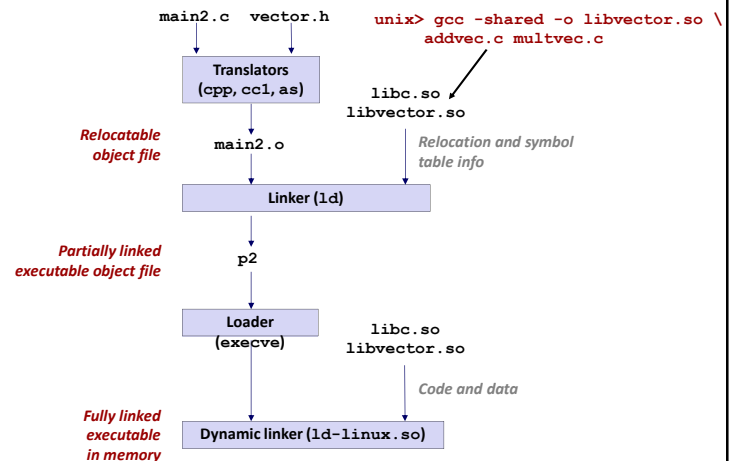
Shared Libraries

- **Static libraries have the following disadvantages:**
 - Duplication in the stored executables (every function need std libc)
 - Duplication in the running executables
 - Minor bug fixes of system libraries require each application to explicitly relink
- **Modern solution: Shared Libraries**
 - Object files that contain code and data that are loaded and linked into an application *dynamically*, at either *load-time* or *run-time*
 - Also called: dynamic link libraries, DLLs, .so files

Shared Libraries (cont.)

- **Dynamic linking can occur when executable is first loaded and run (load-time linking).**
 - Common case for Linux, handled automatically by the dynamic linker (ld-linux.so).
 - Standard C library (libc.so) usually dynamically linked.
- **Dynamic linking can also occur after program has begun (run-time linking).**
 - In Linux, this is done by calls to the `dlopen()` interface.
 - Distributing software.
 - High-performance web servers.
 - Runtime library interpositioning.
- **Shared library routines can be shared by multiple processes.**
 - More on this when we learn about virtual memory

Dynamic Linking at Load-time



Dynamic Linking at Run-time

```
#include <stdio.h>
#include <dlfcn.h>

int x[2] = {1, 2};
int y[2] = {3, 4};
int z[2];

int main()
{
    void *handle;
    void (*addvec)(int *, int *, int *, int);
    char *error;

    /* dynamically load the shared lib that contains addvec() */
    handle = dlopen("./libvector.so", RTLD_LAZY);
    if (!handle) {
        fprintf(stderr, "%s\n", dlerror());
        exit(1);
    }
}
```

Dynamic Linking at Run-time

```
#include <stdio.h>
#include <dlfcn.h>

int x[2] = {1, 2};
int y[2] = {3, 4};
int z[2];

int main()
{
    void *handle;
    void (*addvec)(int *, int *, int *, int);
    char *error;

    /* dynamically load the shared lib that contains addvec() */
    handle = dlopen("./libvector.so", RTLD_LAZY);
    if (!handle) {
        fprintf(stderr, "%s\n", dlerror());
        exit(1);
    }
}
```